

Assignment: "ML with Sklearn"

Abigail Smith

ARS190011

CS 4375.004

Dr. Mazidi

Summary:

This program is for assignment "ML with Sklearn". This program reads in the provided "Auto.csv" and does data exploration. Then the program makes a logistic regression, decision tree, and neural network

```
In [ ]: # Libraries
import pandas as pd
import seaborn as sb
import sklearn as sk
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.metrics import confusion_matrix, classification_report, mean_squared_e
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn import preprocessing
from sklearn.neural_network import MLPRegressor
```

1. Read the Auto data

```
In [ ]: # Prompting user when program starts
print("RUNNING main.py...")

imageHeight = 9 # Set image height for different graphs

# a. Using pandas to read in provided data file
df = pd.read_csv("data//Auto.csv")

# b. Outputting the first few rows
print("\nTop 3 observations from dataframe:")
print(df.head(3))

# c. Outputting the dimensionality
print("\nDimensionality of dataframe:")
print(df.shape)
```

RUNNING main.py...

Top 3 observations from dataframe:

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	\
0	18.0	8	307.0	130	3504	12.0	70.0	
1	15.0	8	350.0	165	3693	11.5	70.0	
2	18.0	8	318.0	150	3436	11.0	70.0	

	origin	name
0	1	chevrolet chevelle malibu
1	1	buick skylark 320
2	1	plymouth satellite

Dimensionality of dataframe:

(392, 9)

2. Data Exploration with Code

```
In [ ]: # a. Using describe() on mpg, weight, and year columns
print("Calling describe() on column \"mpg\"")
print(df["mpg"].describe())
print("\nCalling describe() on column \"weight\"")
print(df["weight"].describe())
print("\nCalling describe() on column \"year\"")
print(df["year"].describe())

# b. Write comments indicating the range and average of each column
# Column      Max      Min      Range      Average
# mpg          46.6       9        37.6       23.45
# weight       5140      1613     3,527     2977.58
# year         82       70       12       76.01
```

Calling describe() on column "mpg"

```
count    392.000000
mean      23.445918
std        7.805007
min        9.000000
25%       17.000000
50%       22.750000
75%       29.000000
max       46.600000
```

Name: mpg, dtype: float64

Calling describe() on column "weight"

```
count    392.000000
mean    2977.584184
std     849.402560
min    1613.000000
25%    2225.250000
50%    2803.500000
75%    3614.750000
max    5140.000000
```

Name: weight, dtype: float64

Calling describe() on column "year"

```
count    390.000000
mean      76.010256
std        3.668093
min       70.000000
25%       73.000000
50%       76.000000
75%       79.000000
max       82.000000
```

Name: year, dtype: float64

3. Explore Data Types

```
In [ ]: # a. Checking the data type of all columns
print("Checking the data types of all columns")
print(df.dtypes)

# b, d. Changing the cylinders column to categorical by cat.codes
print("\nChanging column \"cylinders\" to type category")
print("\tBefore changing:")
print(df["cylinders"].dtypes)
df["cylinders"] = df["cylinders"].astype("category").cat.codes
print("\tAfter changing:")
print(df["cylinders"].dtypes)

# c, d. Changing the origin column to categorical without using cat.codes
print("\nChanging column \"origin\" to type category")
print("\tBefore changing:")
print(df["origin"].dtypes)
df["origin"] = df["origin"].astype("category")
print("\tAfter changing:")
print(df["origin"].dtypes)
```

Checking the data types of all columns

```
mpg            float64
cylinders      int64
displacement   float64
horsepower     int64
weight         int64
acceleration   float64
year           float64
origin         int64
name          object
dtype: object
```

Changing column "cylinders" to type category

Before changing:

```
int64
```

After changing:

```
int8
```

Changing column "origin" to type category

Before changing:

```
int64
```

After changing:

```
category
```

4. Deal with NAs

```
In [ ]: # a. Deleting rows with NAs
print("Count of NAs before removing rows with NAs:")
print(df.isnull().sum())
df = df.dropna()
print("\nCount of NAs after removing rows with NAs:")
print(df.isnull().sum())

# b. Outputting the new dimensions
print("\nDimensionality of dataframe:")
print(df.shape)
```

Count of NAs before removing rows with NAs:

```
mpg          0
cylinders    0
displacement 0
horsepower   0
weight       0
acceleration 1
year         2
origin       0
name         0
dtype: int64
```

Count of NAs after removing rows with NAs:

```
mpg          0
cylinders    0
displacement 0
horsepower   0
weight       0
acceleration 0
year         0
origin       0
name         0
dtype: int64
```

Dimensionality of dataframe:

```
(389, 9)
```

5. Modify Columns

```
In [ ]: # a. Modifying columns
print("Creating new column named \"mpg_high\"")
averageMPG = df["mpg"].mean()
df.insert(0, "mpg_high", 2)

# Writing 0 or 1 to column mpg_high depending on value in column mpg
print("\tBefore updating column \"mpg_high\" values:")
print(df["mpg_high"])

for i, j in df.iterrows():
    if(j["mpg"] >= averageMPG):
        df.at[i, "mpg_high"] = 1
    else:
        df.at[i, "mpg_high"] = 0

print("\tAfter updating column \"mpg_high\" values:")
print(df["mpg_high"])

# b. Dropping column mpg and column name
print("\nDrop columns \"mpg\" and \"name\"")
print("\tAfter dropping column \"mpg\":")
df = df.drop(columns=["mpg", "name"])

# c. Outputting the first few rows of the modified data frame
print("\nTop 5 observations:")
print(df.head(5))
```

Creating new column named "mpg_high"

Before updating column "mpg_high" values:

```
0      2
1      2
2      2
3      2
6      2
```

..

```
387    2
388    2
389    2
390    2
391    2
```

Name: mpg_high, Length: 389, dtype: int64

After updating column "mpg_high" values:

```
0      0
1      0
2      0
3      0
6      0
```

..

```
387    1
388    1
389    1
390    1
391    1
```

Name: mpg_high, Length: 389, dtype: int64

Drop columns "mpg" and "name"

After dropping column "mpg":

Top 5 observations:

	mpg_high	cylinders	displacement	horsepower	weight	acceleration	year	\
0	0	4	307.0	130	3504	12.0	70.0	
1	0	4	350.0	165	3693	11.5	70.0	
2	0	4	318.0	150	3436	11.0	70.0	
3	0	4	304.0	150	3433	12.0	70.0	
6	0	4	454.0	220	4354	9.0	70.0	

origin

```
0      1
1      1
2      1
3      1
6      1
```

6. Data Exploration with Graphs

d. For each graph, write a comment indicating one thing you learned about the data from the graph.


catplot (mpg_high): This plot shows the distribution of the values in column "mpg_high". This graph shows that the new column "mpg_high" has only 0 and 1 values. We can also see there's nearly the same amount of 0 and 1 in column "mpg_high".

catplot (horsepower, weight, mpg_high): This plot shows the distribution of horsepower by weight where each dot's color indicates whether the observation was labeled with a high mpg or a low mpg. From this graph, we can see that majority of high mpgs have a lower horsepower and a lower weight. In contrast, we can also see that majority of low mpgs have a higher horsepower and a higher weight.

boxplot: This plot shows the boxplot of mpg_high to weight. The minimum weight for a lower mpg is higher than the minimum weight for a higher mpg. The median weight for a lower mpg is higher than the median weight for a higher mpg. The maximum weight for a lower mpg is higher than the maximum weight for a higher mpg. Additionally, a lower mpg has a wider range than the weight for a higher mpg. Finally, there are more outliers for a higher mpg than a lower mpg.


```
In [ ]: # a. Using seaborn catplot on column mpg_high
        sb.catplot(df, x = "mpg_high", height = imageHeight)
```

```
Out[ ]: <seaborn.axisgrid.FacetGrid at 0x285c9613a30>
```



```
In [ ]: # b. Replotting with column "horsepower" on x-axis, column "weight" on y-axis,  
# and setting hue set to column "mpg_high"  
sb.catplot(data = df, x = "horsepower", y = "weight", hue = "mpg_high", height = im
```

```
Out[ ]: <seaborn.axisgrid.FacetGrid at 0x285ca890d00>
```

```
In [ ]: # pg,cylinders,displacement,horsepower,weight,acceleration,year
sb.catplot(data = df, x = "horsepower", y = "cylinders", hue = "mpg_high", height = 4)
sb.catplot(data = df, x = "horsepower", y = "displacement", hue = "mpg_high", height = 4)
sb.catplot(data = df, x = "horsepower", y = "weight", hue = "mpg_high", height = 4)
sb.catplot(data = df, x = "horsepower", y = "acceleration", hue = "mpg_high", height = 4)
sb.catplot(data = df, x = "horsepower", y = "year", hue = "mpg_high", height = 4)


sb.catplot(data = df, x = "cylinders", y = "displacement", hue = "mpg_high", height = 4)
sb.catplot(data = df, x = "cylinders", y = "weight", hue = "mpg_high", height = 4)
sb.catplot(data = df, x = "cylinders", y = "acceleration", hue = "mpg_high", height = 4)
sb.catplot(data = df, x = "cylinders", y = "year", hue = "mpg_high", height = 4)


sb.catplot(data = df, x = "displacement", y = "weight", hue = "mpg_high", height = 4)
sb.catplot(data = df, x = "displacement", y = "acceleration", hue = "mpg_high", height = 4)
sb.catplot(data = df, x = "displacement", y = "year", hue = "mpg_high", height = 4)


sb.catplot(data = df, x = "weight", y = "acceleration", hue = "mpg_high", height = 4)
sb.catplot(data = df, x = "weight", y = "year", hue = "mpg_high", height = 4)


sb.catplot(data = df, x = "acceleration", y = "year", hue = "mpg_high", height = 4)
```


```
Out[ ]: <seaborn.axisgrid.FacetGrid at 0x285d405ac80>
```










































```
In [ ]: # c. Using seaborn to boxplot with column "mpg_high" on x-axis and weight on y-axis
sb.boxplot(df, x = "mpg_high", y = "weight")
```

```
Out[ ]: <AxesSubplot: xlabel='mpg_high', ylabel='weight'>
```

7. Train/Test Split

```
In [ ]: # a, b, c. Creating 80/20 split and using seed 1234 to get same results each run
X = df.drop(columns = ["mpg_high"])
X_train, X_test, Y_train, Y_test = train_test_split(X, df["mpg_high"], test_size =

# d. Output the dimensions of train and test
print("Showing size of test and train subsets:")
print("\tSize of test: ", X_test.shape)
print("\tSize of train: ", X_train.shape)
```

Showing size of test and train subsets:

```
Size of test: (78, 7)
Size of train: (311, 7)
```

8. Logistic Regression

```
In [ ]: # a, b. Making Logistic regression model
logisticRegression = LogisticRegression(solver = "lbfgs", max_iter=100000)
logisticRegression.fit(X_train, Y_train)
logisticRegression.score(X_train, Y_train)

# Making prediction on Logistic regression model
pred = logisticRegression.predict(X_test)

# Finding accuracy score, precision score, recall score, and f1 score
accuracy = accuracy_score(Y_test, pred)
precision = precision_score(Y_test, pred)
recall = recall_score(Y_test, pred)
```

```

f1 = f1_score(Y_test, pred)
score = logisticRegression.score(X_test, Y_test)

print("Logistic Regression:")
print("\tAccuracy score: ", accuracy)
print("\tPrecision score: ", precision)
print("\tRecall score: ", recall)
print("\tf1 score: ", f1)
print("\tScore: ", score)

# Making confusion matrix
confusion_matrix(Y_test, pred)

# c. Printing classification report
print("\n\tClassification Report:")
print(classification_report(Y_test, pred))

```

Logistic Regression:

```

Accuracy score: 0.8974358974358975
Precision score: 0.7777777777777778
Recall score: 1.0
f1 score: 0.8750000000000001
Score: 0.8974358974358975

```

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.84	0.91	50
1	0.78	1.00	0.88	28
accuracy			0.90	78
macro avg	0.89	0.92	0.89	78
weighted avg	0.92	0.90	0.90	78

9. Decision Tree

```

In [ ]: # a. Training a decision tree
decisionTree = DecisionTreeClassifier()
decisionTree.fit(X_train, Y_train)

# b. Making prediction on decision tree model
pred = decisionTree.predict(X_test)

# Finding accuracy score, precision score, recall score, and f1 score
accuracy = accuracy_score(Y_test, pred)
precision = precision_score(Y_test, pred)
recall = recall_score(Y_test, pred)
f1 = f1_score(Y_test, pred)
score = decisionTree.score(X_test, Y_test)

print("Decision Tree:")
print("\tAccuracy score: ", accuracy)
print("\tPrecision score: ", precision)
print("\tRecall score: ", recall)
print("\tf1 score: ", f1)

```

```

print("\tScore: ", score)

# Making confusion matrix
confusion_matrix(Y_test, pred)

# c. Printing classification report
print("\n\tClassification Report:")
print(classification_report(Y_test, pred))

```

Decision Tree:

```

Accuracy score: 0.8717948717948718
Precision score: 0.8
Recall score: 0.8571428571428571
f1 score: 0.8275862068965518
Score: 0.8717948717948718

```

Classification Report:

	precision	recall	f1-score	support
0	0.92	0.88	0.90	50
1	0.80	0.86	0.83	28
accuracy			0.87	78
macro avg	0.86	0.87	0.86	78
weighted avg	0.87	0.87	0.87	78

10. Neural Network

e. Compare the two models and why you think the performance was same/different In neural networks, each node in the hidden layers learns a different function from the inputs. With too few nodes, underfitting can occur and vice-versa. The differences in the output occur due to the dataset size combined with the different number of layers and nodes influencing the model's chances of underfitting or overfitting. As seen when the `hidden_layer_size` was set to (6,3) and when the hidden layer size was set to (15, 10, 5), the MSE, correlation, and score are all impacted by the number of hidden layers and the number of nodes within each layer. From these outputs, it would be a fair assumption to say that as the number of layers and nodes increase from (6, 3), the worse the neural network for this dataset performs, as seen by the stark difference between the MSE, correlation, and score between the two models.

Though we were only asked to test two neural networks, I went ahead and created three -- simply because I was interested in seeing the outputs for much larger number of nodes in each layer. As seen in the third output, having a `hidden_layer_size` set to (100, 50) produced a lower MSE and higher score and correlation than the second model, but a worse correlation, score, and MSE to the first model.

By the structure of neural networks, changing the number of layers and nodes directly impact the output of the neural network models.

```
In [ ]: # a. train a neural network, choosing a network topology of your choice
        scaler = preprocessing.StandardScaler().fit(X_train)

        X_train = scaler.transform(X_train)
        X_test = scaler.transform(X_test)

        # b. First neural network model with 6 nodes in the first hidden layer and 3 nodes
        # hidden layer
        regr = MLPRegressor(hidden_layer_sizes = (6, 3), max_iter = 5000, random_state = 12)
        regr.fit(X_train, Y_train)
        Y_pred = regr.predict(X_test)

        mse = mean_squared_error(Y_test, Y_pred)
        correlation = r2_score(Y_test, Y_pred)
        score = regr.score(X_test, Y_test)

        print("1. hidden_layer_sizes = (6, 3)")
        print("\tMSE: ", mse)
        print("\tCorrelation: ", correlation)
        print("\tScore: ", score)

        # c, d. Second neural network model with 15 nodes in the first hidden layer, 10 nodes
        # hidden layer, and 5 nodes in the third hidden layer
        regr = MLPRegressor(hidden_layer_sizes = (15, 10, 5), max_iter = 5000, random_state = 12)
        regr.fit(X_train, Y_train)
        Y_pred = regr.predict(X_test)

        mse = mean_squared_error(Y_test, Y_pred)
        correlation = r2_score(Y_test, Y_pred)
        score = regr.score(X_test, Y_test)

        print("\n2. hidden_layer_sizes = (15, 10, 5)")
        print("\tMSE: ", mse)
        print("\tCorrelation: ", correlation)
        print("\tScore: ", score)

        # Third neural network model with 7 nodes in the first hidden layer and 4 nodes in
        # hidden layer
        regr = MLPRegressor(hidden_layer_sizes = (7, 4), max_iter = 5000, random_state = 12)
        regr.fit(X_train, Y_train)
        Y_pred = regr.predict(X_test)

        mse = mean_squared_error(Y_test, Y_pred)
        correlation = r2_score(Y_test, Y_pred)
        score = regr.score(X_test, Y_test)

        print("\n3. hidden_layer_sizes = (100, 50)")
        print("\tMSE: ", mse)
        print("\tCorrelation: ", correlation)
        print("\tScore: ", score)
```

1. hidden_layer_sizes = (6, 3)
MSE: 0.08188287460306254
Correlation: 0.6441604220821195
Score: 0.6441604220821195
2. hidden_layer_sizes = (15, 10, 5)
MSE: 0.11553113107870458
Correlation: 0.4979347132265437
Score: 0.4979347132265437
3. hidden_layer_sizes = (100, 50)
MSE: 0.08718585115782261
Correlation: 0.6211152011112908
Score: 0.6211152011112908